

Seeing Is No Longer Believing

Deepfakes, scam calls, and the collapse of proof

CEP146 — Essential Tooling for Programmers | Group 3

Three calls that never happened

Same technology, three different victims

1

A finance worker

Joins a video call with his CFO and colleagues. Every face on the call is generated. He wires out \$25.6M.

2

A retail investor

Watches the CEO of a stock exchange give him a hot tip on WhatsApp. The CEO never recorded it.

3

A parent


Picks up the phone and hears their own child's voice asking for money. Three seconds of audio was enough to clone it.

0.1%

of people could reliably
tell real from fake
in a 2025 study

It is not magic — it is a copy machine for people

What the technology actually does



The hard part used to be the compute. That bill is now a few dollars an hour, rented by the minute.

**The cloud removed the barrier
between an idea and the world.**

It removed the same barrier for a lie.

Where it actually lands

Fraud is only the first layer



Your phone

Cloned voices turn a scam call into a family emergency. Social media is the single largest origin of reported losses.



Your feed

A fake image of a collapsed bridge after the December 2025 UK earthquake was convincing enough to cancel trains.



Your money

In January 2026 the Bombay Stock Exchange warned investors about a deepfake of its own CEO. It resurfaced by March.

The damage is not always a transfer. Sometimes it is just doubt — and doubt is cheaper to manufacture.

This is now a line item, not a novelty

Reported figures, 2025 into 2026

\$893M

in AI-related fraud losses logged by the FBI in 2025 — the first year it counted them separately

\$1.65B

of the \$2.19B in deepfake fraud recorded since 2019 was lost in 2025 alone

1 in 5

biometric fraud attempts is now a deepfake, across a billion identity checks

Read the numbers carefully. Vendor blogs quote surges of 900% to 3,892% that contradict each other. We used the FBI and a study built on public incident records — and even those are undercounts, because most victims never report.

“I would be able to tell”

Almost nobody can, and the tools are not much better

- Humans separate real from synthetic at roughly coin-flip accuracy — while most people are confident they would do better
- Detection tools that hit ~96% in the lab lose 45–50% of that accuracy on real-world footage
- Every fix that gets published becomes training data for the next generation of fakes

~50%

human accuracy on
high-quality fakes.
A coin does the same job.

So the defence cannot be “look closer.” It has to be a process.

When the target is a person, not a bank

The harm with no invoice attached

The most common use of this technology has never been financial. It is fake intimate imagery of real people — overwhelmingly women, and increasingly minors — made from ordinary photos anyone can download.

As of 19 May 2026, US platforms must give victims a way to report that content and remove it, and known copies, within 48 hours. The FTC wrote to Meta, Google, TikTok, X, Reddit, Snap and others before the deadline to say it would enforce.

48 hours

to remove non-consensual
imagery after a valid request,
under the TAKE IT DOWN Act
(penalties up to \$53,088
per violation)

What is actually being done

Law, provenance, and one habit that costs nothing

Label it

Since 2 August 2026, the EU AI Act requires deepfakes and AI-generated content to be disclosed. Fines reach €15M or 3% of global turnover.

Sign it

Provenance standards attach a machine-readable trail to real footage — proving what is genuine rather than hunting for what is fake.

Verify it

Hang up and call back on a number you already had. Agree a family code word. It defeats a \$25M attack for free.

Regulation sets the floor. The habit is what protects you tonight.

Discussion question

If a video can no longer prove anything, what do we accept as proof instead?

Thanks for watching — Group 3